



Introduction to Machine Learning

NYU K12 STEM Education: Machine Learning

Department of Electrical and Computer Engineering,
NYU Tandon School of Engineering
Brooklyn, New York

- ▶ [Course Website](#)
- ▶ Instructors:



Rugved Mhatre
rugved.mhatre@nyu.edu



Akshath Mahajan
akshathmahajan@nyu.edu

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
4. Data Augmentation
5. Normalization
6. Dropout
7. Transfer Learning

Review
●●

Working with Images
○○○○

Convolution Neural Networks
○○○○○○○

Data Augmentation
○○○○○○○

Normalization
○○○○

Dropout
○○○

Transfer Learning
○○○

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
4. Data Augmentation
5. Normalization
6. Dropout
7. Transfer Learning

Grayscale Images

- Images are stored as arrays of quantized numbers in computers

Grayscale Images

- ▶ Images are stored as arrays of quantized numbers in computers
- ▶ 2D matrices with each entry specifying the intensity (brightness) of a pixel

Grayscale Images

- ▶ Images are stored as arrays of quantized numbers in computers
- ▶ 2D matrices with each entry specifying the intensity (brightness) of a pixel
- ▶ Pixel values range from 0 to 255, 0 being the darkest, 255 being the brightest

```
[[255 255 255]  
 [255  0 255]  
 [255 255 255]]
```




Figure 1: A 3x3 Grayscale Image

Color Images

- ▶ Color (RGB) images have an extra dimension for color (3D array)
- ▶ Imagine three 2D matrices stacked together
- ▶ Each 2D matrix specifies the amount of color for Red, Green, and Blue at each pixel




Figure 2: RGB Images

Color Images

- ▶ Color (RGB) images have an extra dimension for color (3D array)
- ▶ Imagine three 2D matrices stacked together
- ▶ Each 2D matrix specifies the amount of color for Red, Green, and Blue at each pixel




Figure 2: RGB Images

- ▶ Shape - (1050, 700, 3)

Images and Neural Networks

- How to feed Images in a Fully Connected Network?

Images and Neural Networks

- ▶ How to feed Images in a Fully Connected Network?
- ▶ Flatten the image!




Figure 3: Flattening an Image

Images and Neural Networks

- ▶ How to feed Images in a Fully Connected Network?
- ▶ Flatten the image!




Figure 3: Flattening an Image

- ▶ Does this make sense? Is this how we see images?

Images and Neural Networks

- ▶ How to feed Images in a Fully Connected Network?
- ▶ Flatten the image!




Figure 3: Flattening an Image

- ▶ Does this make sense? Is this how we see images?
 - ▶ No consideration for spatial positions!!
 - ▶ How many input neurons for 1024x1024 image?
 - ▶ What about slightly rotated photographs?

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
4. Data Augmentation
5. Normalization
6. Dropout
7. Transfer Learning

The Convolution Operation

- ▶ All these problems are solved by Convolutions!
- ▶ Convolution operation is applied on an image matrix X with a kernel W

$$Z = X \circledast W$$

The Convolution Operation

- ▶ All these problems are solved by Convolutions!
- ▶ Convolution operation is applied on an image matrix X with a kernel W

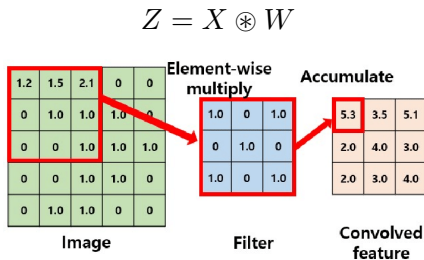


Figure 4: Convolution Operation

The Convolution Operation

► Let's see some visualizations!

Why Convolution?

- ▶ With convolution, each output pixel depends on only the neighboring pixels in the input

Why Convolution?

- ▶ With convolution, each output pixel depends on only the neighboring pixels in the input
- ▶ This allows us to learn the positional relationship between pixels

Why Convolution?

- ▶ With convolution, each output pixel depends on only the neighboring pixels in the input
- ▶ This allows us to learn the positional relationship between pixels
- ▶ Different kernels capture different features from the image

Convolution for Multiple Channels

- There is a single kernel for each channel

Convolution for Multiple Channels

- ▶ There is a single kernel for each channel
- ▶ Each kernel performs a 2D convolution on its respective channel

Convolution for Multiple Channels

- ▶ There is a single kernel for each channel
- ▶ Each kernel performs a 2D convolution on its respective channel
- ▶ The results are then summed

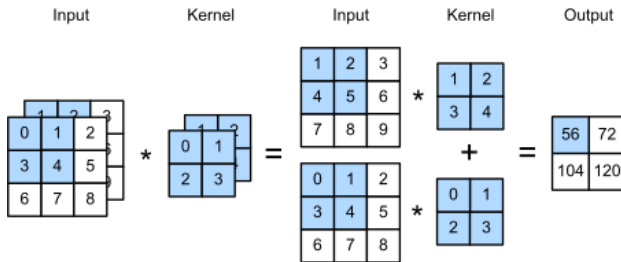


Figure 5: Convolution Across Channels

Source dl2.ai

Max Pooling

- ▶ It is a down-sampling technique in Convolutional Neural Networks
- ▶ Reduces the dimensions of intermediate network results

Max Pooling

- ▶ It is a down-sampling technique in Convolutional Neural Networks
- ▶ Reduces the dimensions of intermediate network results
- ▶ It provides "translational invariance". Why?

Max Pooling

- ▶ It is a down-sampling technique in Convolutional Neural Networks
- ▶ Reduces the dimensions of intermediate network results
- ▶ It provides "translational invariance". Why?
 - ▶ Most prominent feature in every local region is preserved
 - ▶ Focuses on the presence of features rather than their precise location

Max Pooling

- Let's see an example!




Figure 6: Max Pooling Example

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
- 4. Data Augmentation**
5. Normalization
6. Dropout
7. Transfer Learning

Scarcity of training data

- ▶ Large-scale deep learning models are extremely data hungry

Scarcity of training data

- ▶ Large-scale deep learning models are extremely data hungry
- ▶ We don't always have enough data to train the model
- ▶ Labelling data is expensive and time-consuming

Scarcity of training data

- ▶ Large-scale deep learning models are extremely data hungry
- ▶ We don't always have enough data to train the model
- ▶ Labelling data is expensive and time-consuming
- ▶ What can we do now?

Scarcity of training data

- ▶ Large-scale deep learning models are extremely data hungry
- ▶ We don't always have enough data to train the model
- ▶ Labelling data is expensive and time-consuming
- ▶ What can we do now?
- ▶ Create new images!

Data Augmentation

- ▶ Key Idea: Augment existing images from the original dataset

Data Augmentation

- ▶ Key Idea: Augment existing images from the original dataset
- ▶ Similar enough to contain the same Subject as the original
- ▶ Different enough to prove meaningful for training

Data Augmentation

- ▶ Key Idea: Augment existing images from the original dataset
- ▶ Similar enough to contain the same Subject as the original
- ▶ Different enough to prove meaningful for training
- ▶ Let's look at some techniques for Data Augmentation

Mirroring




Figure 7: Mirroring

Rotation and Translation




Figure 8: Rotation and Translation

Random Cropping

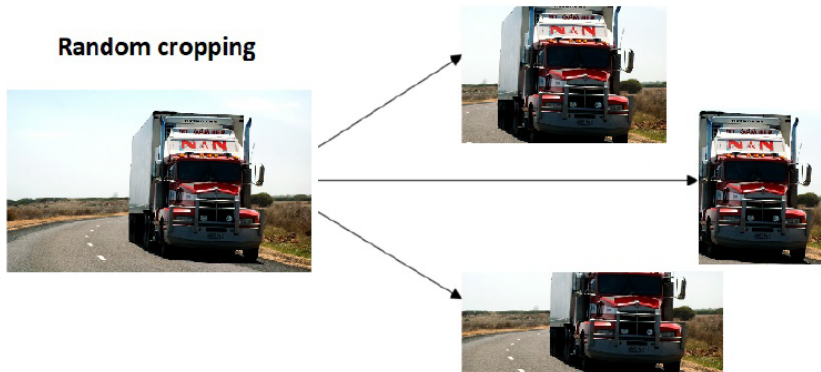


Figure 9: Random Cropping

Color Shifting




Figure 10: Color Shifting

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
4. Data Augmentation
- 5. Normalization**
6. Dropout
7. Transfer Learning

Data Normalization

- Consider the dataset $(x_i, y_i) \forall i \in \{1, 2, 3, \dots, N\}$

Data Normalization

- ▶ Consider the dataset $(x_i, y_i) \forall i \in \{1, 2, 3, \dots, N\}$
- ▶ Mean $\bar{x} = \frac{1}{N} \sum x_i$
- ▶ Variance $\sigma^2 = \frac{1}{N} \sum (x_i - \bar{x})^2$

Data Normalization

- ▶ Consider the dataset $(x_i, y_i) \forall i \in \{1, 2, 3, \dots, N\}$
- ▶ Mean $\bar{x} = \frac{1}{N} \sum x_i$
- ▶ Variance $\sigma^2 = \frac{1}{N} \sum (x_i - \bar{x})^2$
- ▶ Normalization: Replace each x_i with x'_i , where:

$$x'_i = \frac{x_i - \bar{x}}{\sigma}$$

- ▶ The new dataset will have a mean of 0 and a variance of 1

Data Normalization

- ▶ Consider a single weight w and bias b
- ▶ The contours in the plot represents the value of the loss function for the given w and b




Figure 11: Unnormalized vs Normalized Descent
Source: [TowardsDataScience](#)

Batch Normalization

- We can normalize inputs to the network. Why not do that to the intermediate layer outputs

Batch Normalization

- ▶ We can normalize inputs to the network. Why not do that to the intermediate layer outputs
- ▶ Batch Normalization involves normalizing the inputs to each layer within each mini-batch

Batch Normalization

- ▶ We can normalize inputs to the network. Why not do that to the intermediate layer outputs
- ▶ Batch Normalization involves normalizing the inputs to each layer within each mini-batch
- ▶ Batch normalization is applied before activation

```
model = models.Sequential()

model.add(layers.Conv2D(64,(3,3)))
model.add(layers.BatchNormalization())
model.add(layers.Activation('relu'))

model.add(layers.Flatten())
model.add(layers.Dense(64))
model.add(layers.BatchNormalization())
model.add(layers.Activation('relu'))

model.add(layers.Dense(1, activation =
'sigmoid'))
```

Figure 12: Batch Normalization

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
4. Data Augmentation
5. Normalization
- 6. Dropout**
7. Transfer Learning

Dropout

- ▶ This technique is patented by Google

Dropout

- ▶ This technique is patented by Google
- ▶ Randomly disable neurons and their connections between each other




Figure 13: Dropout

Dropout

- ▶ This technique is patented by Google
- ▶ Randomly disable neurons and their connections between each other
- ▶ Without dropout, neurons can become too reliant on the outputs of specific other neurons, leading to overfitting




Figure 13: Dropout

Dropout

- This is the same as using a neural network with the same amount of layers but less neurons per layer.




Figure 14: Dropout

Dropout

- ▶ This is the same as using a neural network with the same amount of layers but less neurons per layer.
- ▶ The more neurons the more powerful the neural network is, and the more likely it is to overfit.




Figure 14: Dropout

Dropout

- ▶ This is the same as using a neural network with the same amount of layers but less neurons per layer.
- ▶ The more neurons the more powerful the neural network is, and the more likely it is to overfit.
- ▶ This also means that the model can not rely on any single feature, therefore would need to spread out the weights




Figure 14: Dropout

Outline

1. Review
2. Working with Images
3. Convolution Neural Networks
4. Data Augmentation
5. Normalization
6. Dropout
- 7. Transfer Learning**

Transfer Learning

- ▶ Key Idea: "Standing on the shoulder of Giants"

Transfer Learning

- ▶ Key Idea: "Standing on the shoulder of Giants"
- ▶ Training large computer vision models requires extensive hyperparameter search and multiple GPU running for weeks!

Transfer Learning

- ▶ Key Idea: "Standing on the shoulder of Giants"
- ▶ Training large computer vision models requires extensive hyperparameter search and multiple GPU running for weeks!
- ▶ Solution: Transfer Learning




Figure 15: Transfer Learning

Source: [Oreilly](#)

Transfer Learning

- Researchers now open-source their model weights, which can be a great initialization point for your applications

Transfer Learning

- ▶ Researchers now open-source their model weights, which can be a great initialization point for your applications
- ▶ Often in practice, people preserve the feature extractor and re-train the classification head

Transfer Learning

- ▶ Researchers now open-source their model weights, which can be a great initialization point for your applications
- ▶ Often in practice, people preserve the feature extractor and re-train the classification head
- ▶ Freeze the early layers and replace the last few to match your needs. Only train the replaced layers

Transfer Learning

- ▶ Researchers now open-source their model weights, which can be a great initialization point for your applications
- ▶ Often in practice, people preserve the feature extractor and re-train the classification head
- ▶ Freeze the early layers and replace the last few to match your needs. Only train the replaced layers
- ▶ This is similar to transferring the knowledge from one network to another, thus the name transfer learning.